



# CHƯƠNG 6: ĐỒNG BỘ HÓA

TS. TRẦN HẢI ANH

Tham khảo bài giảng của PGS. TS. Hà Quốc Trung

# Contents



- Clock synchronization
- Logical clock
- Mutual exclusion
- Election algorithm

# 1. Clock Synchronization



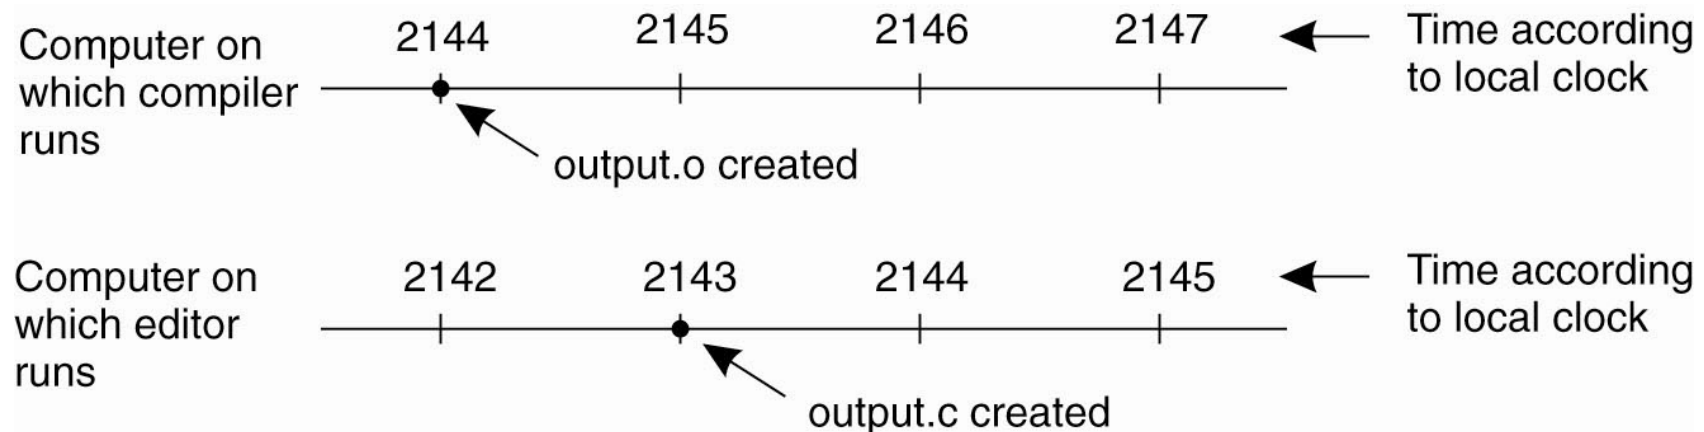
- Notion of synchronization
- Physical Clocks
- Global Positioning System
- Clock Synchronization Algorithms
- Use of Synchronized Clocks

# Synchronization



- How process synchronize
  - ▣ Multiple process to not simultaneously access to the same resources: printers, files
  - ▣ Multiple process are agreed on the ordering of event.
    - Ex: message m1 of P is sent after m2 of Q
- Synchronization based on actual time
- Synchronization by relative ordering

# Clock Synchronization



- When each machine has its own clock, an event that occurred after another event may nevertheless be assigned an earlier time.

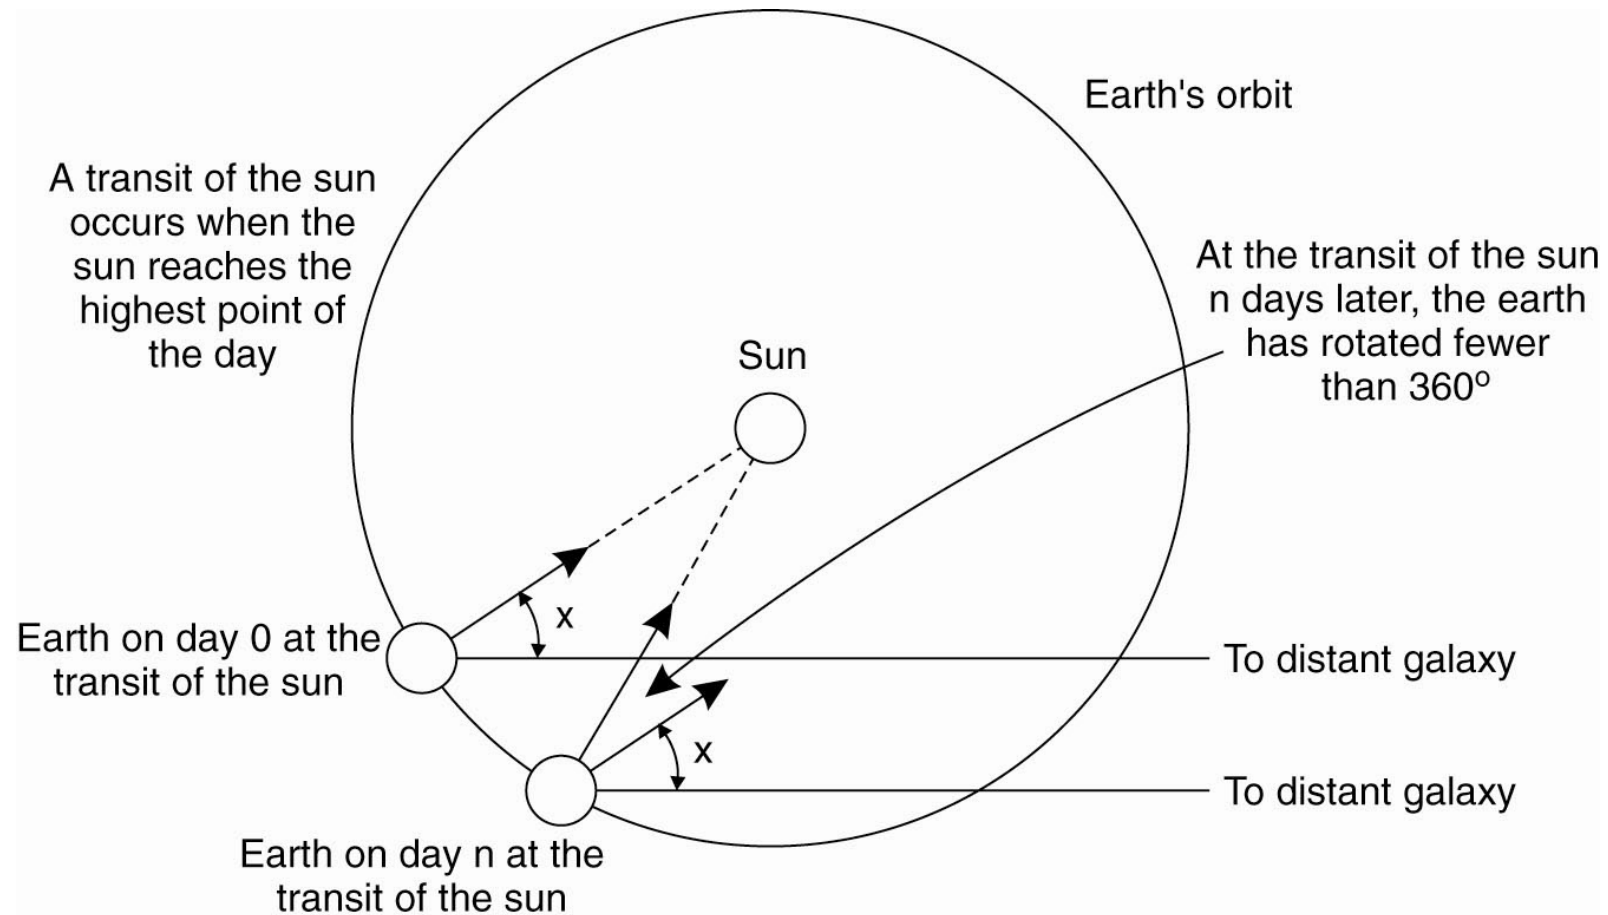
# Physycal Clocks

- ❑ Timer
- ❑ Counter & Holding register
- ❑ Clock tick
- ❑ Problem in distributed systems:
  - ▣ How do we synchronize them with real-world?
  - ▣ How do we synchronize the clocks with each other?



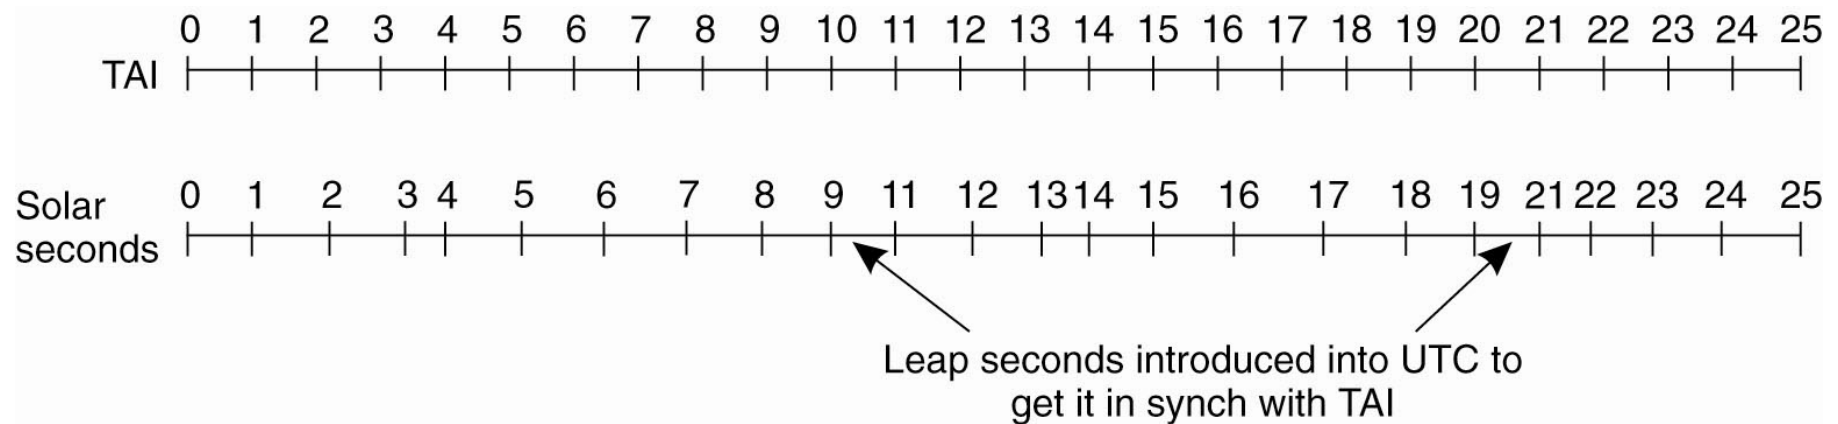
**RTC IC  
(Real Time Clock)**

# Physical Clocks (1)



**Computation of the mean solar day**

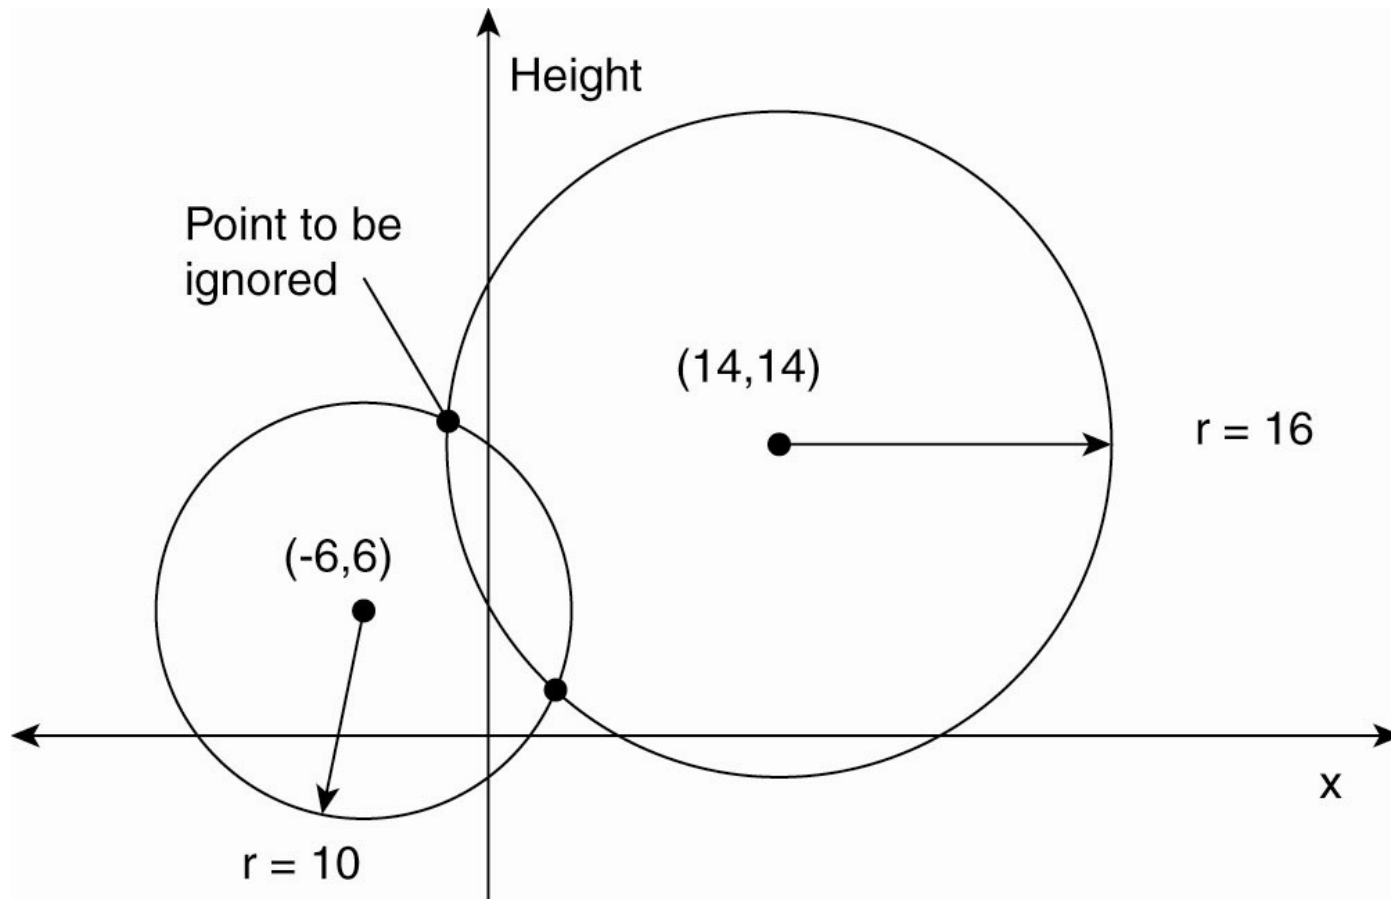
# Physical Clocks (2)



- TAI seconds are of constant length, unlike solar seconds. Leap seconds are introduced when necessary to keep in phase with the sun.
- => UTC (Universal Cordinated Time)



# Global Positioning System (1)



**Computing a position in a two-dimensional space**

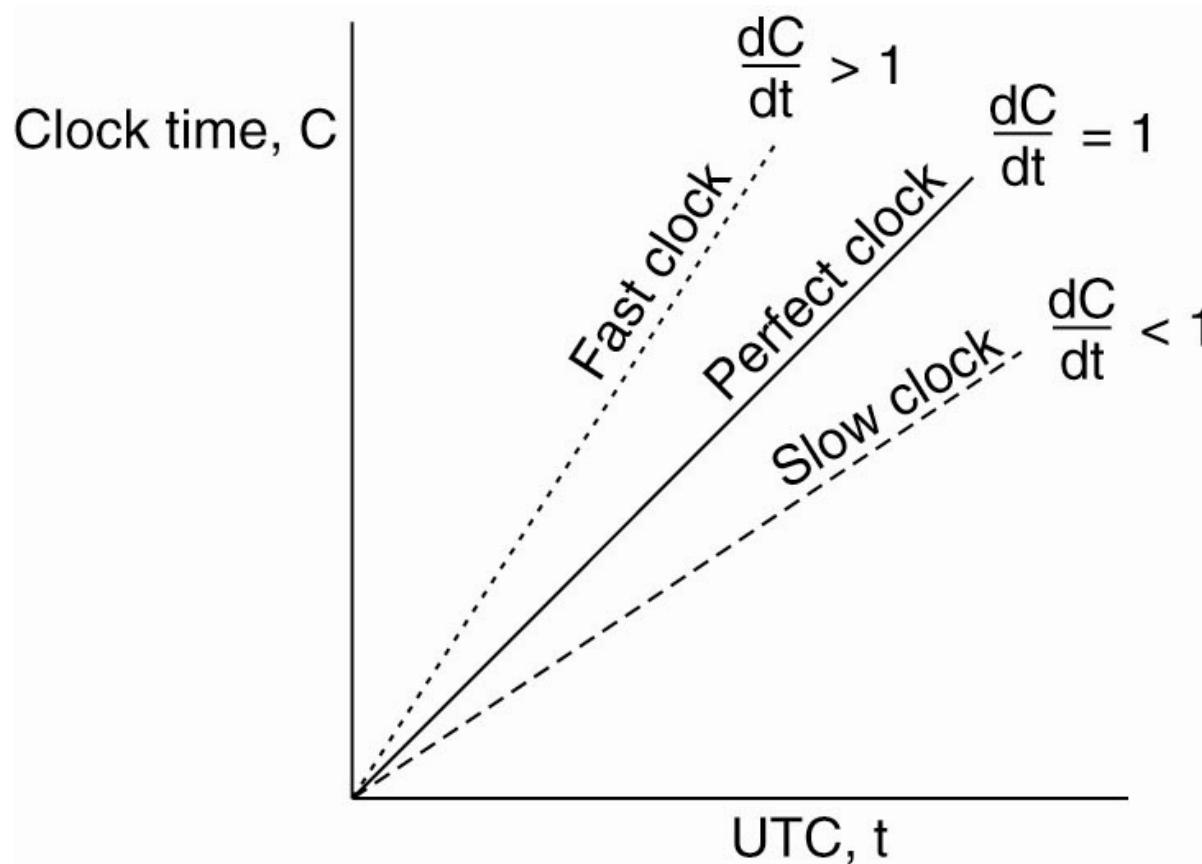
# Global Positioning System (2)



- Real world facts that complicate GPS
  1. It takes a while before data on a satellite's position reaches the receiver.
  2. The receiver's clock is generally not in synch with that of a satellite.

# Clock Synchronization Algorithms

- The relation between clock time and UTC when clocks tick at different rates.

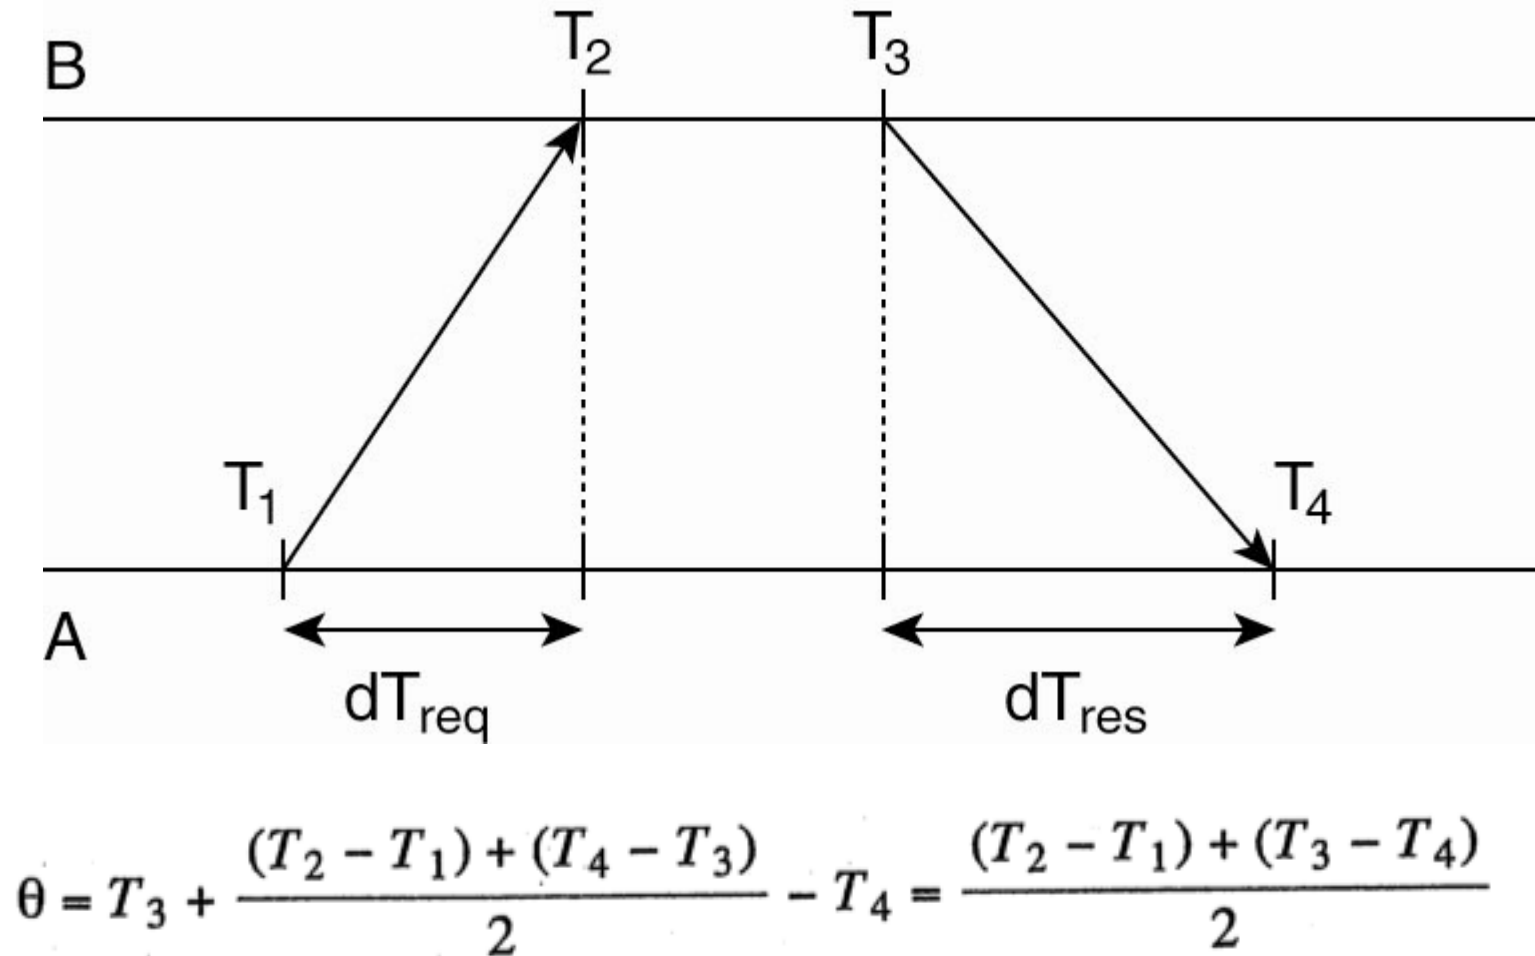


# Clock Synchronization Algorithms



- Network Time Protocol
- Berkeley Algorithm
- Clock Synchronization in Wireless Networks

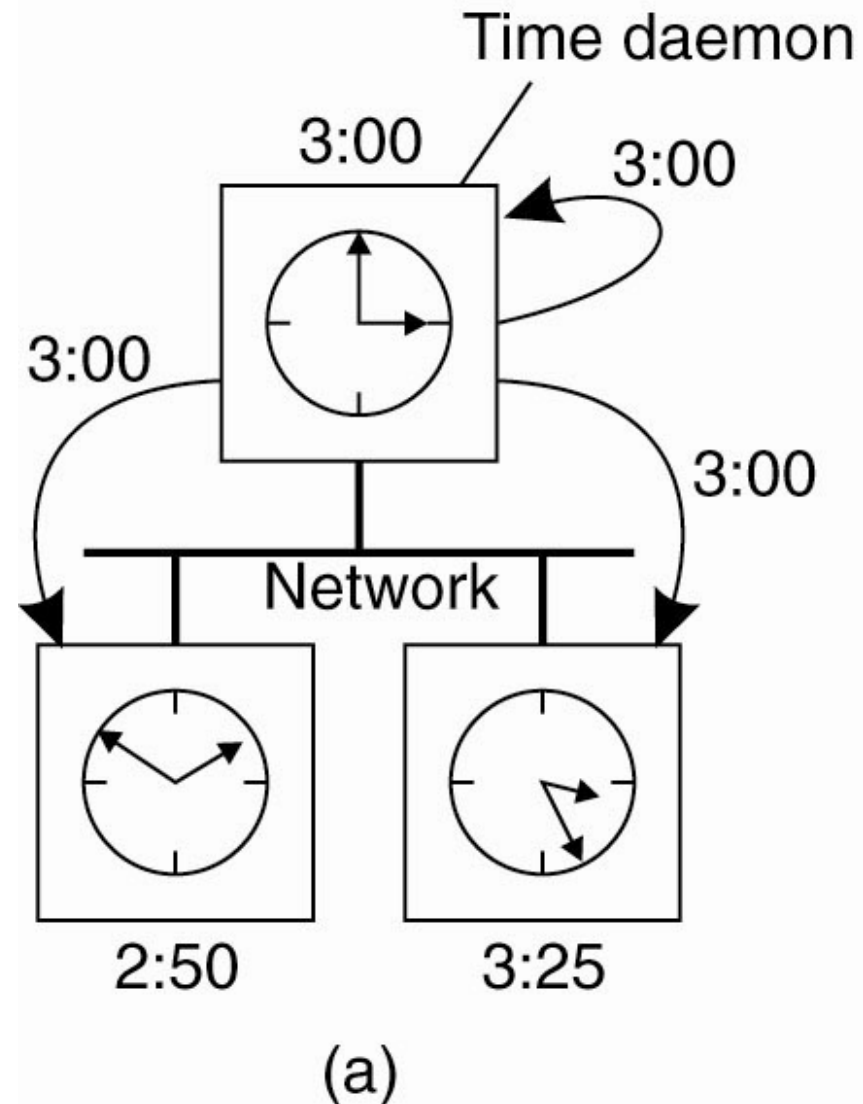
# Network Time Protocol



**Getting the current time from a time server**

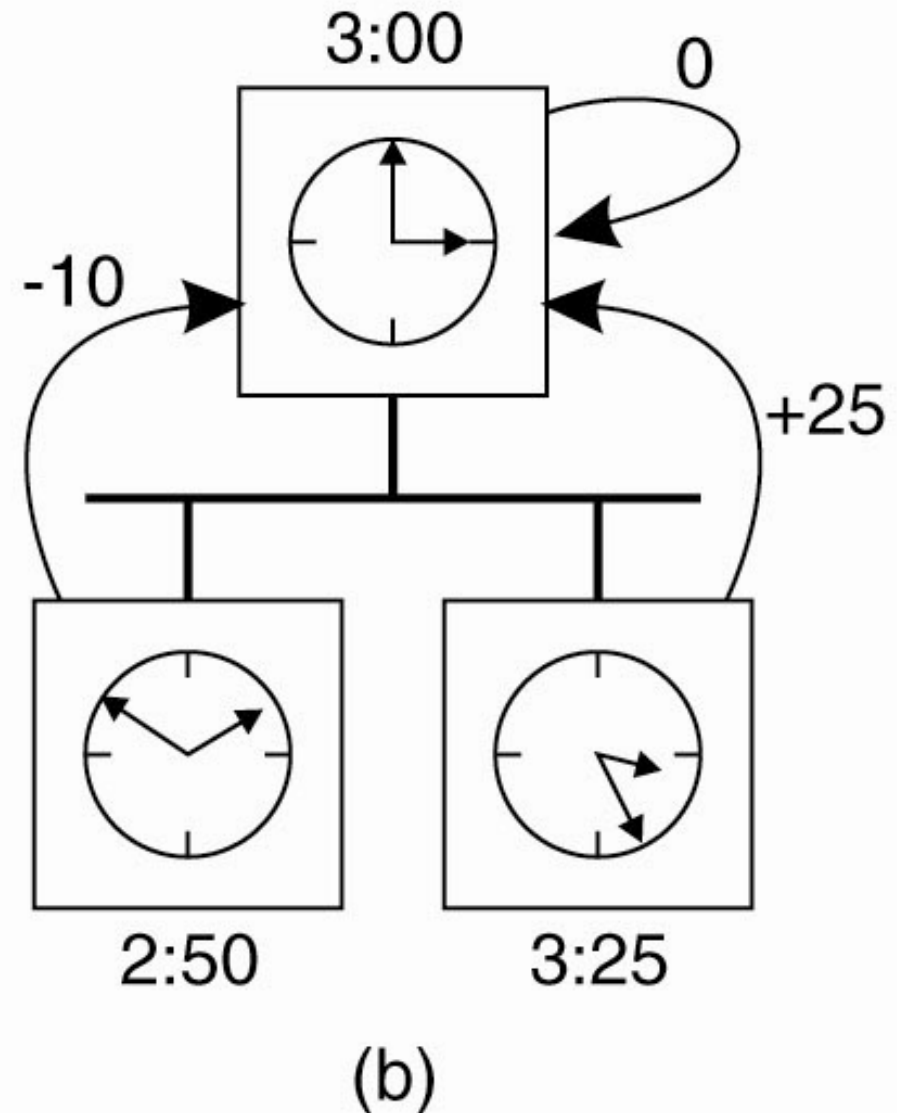
# The Berkeley Algorithm (1)

- The time daemon asks all the other machines for their clock values.



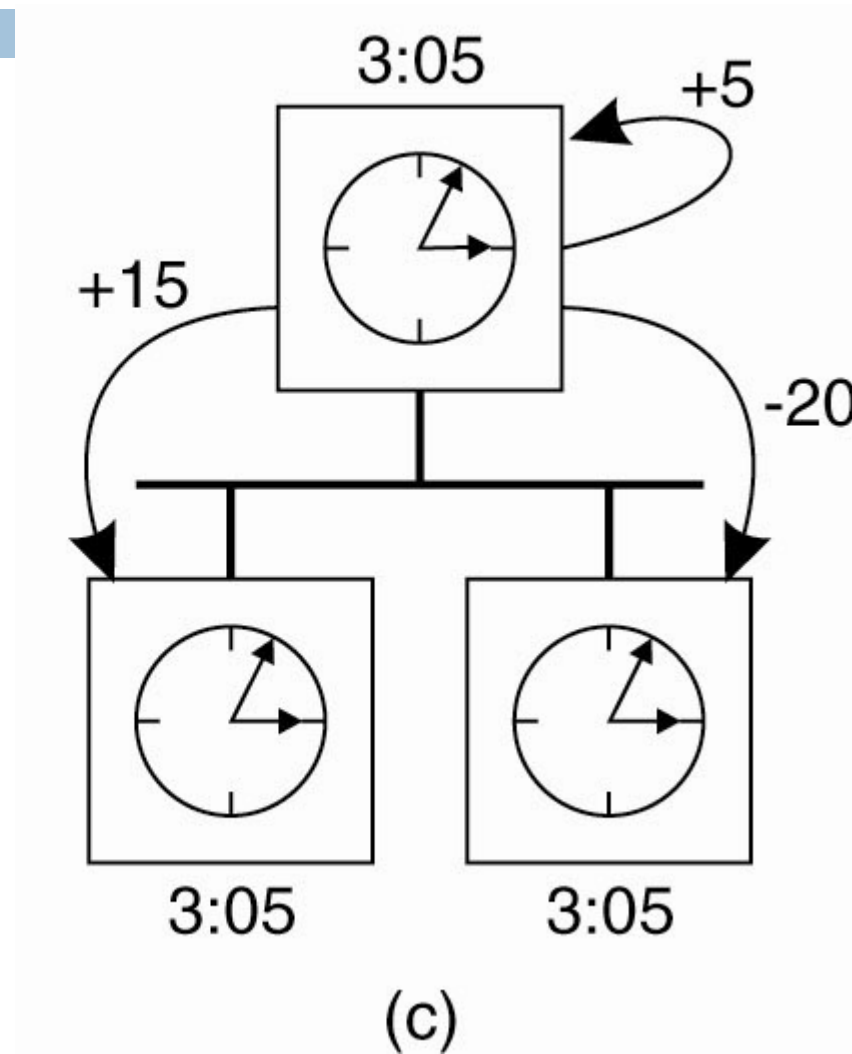
# The Berkeley Algorithm (2)

- The machines answer.



# The Berkeley Algorithm (3)

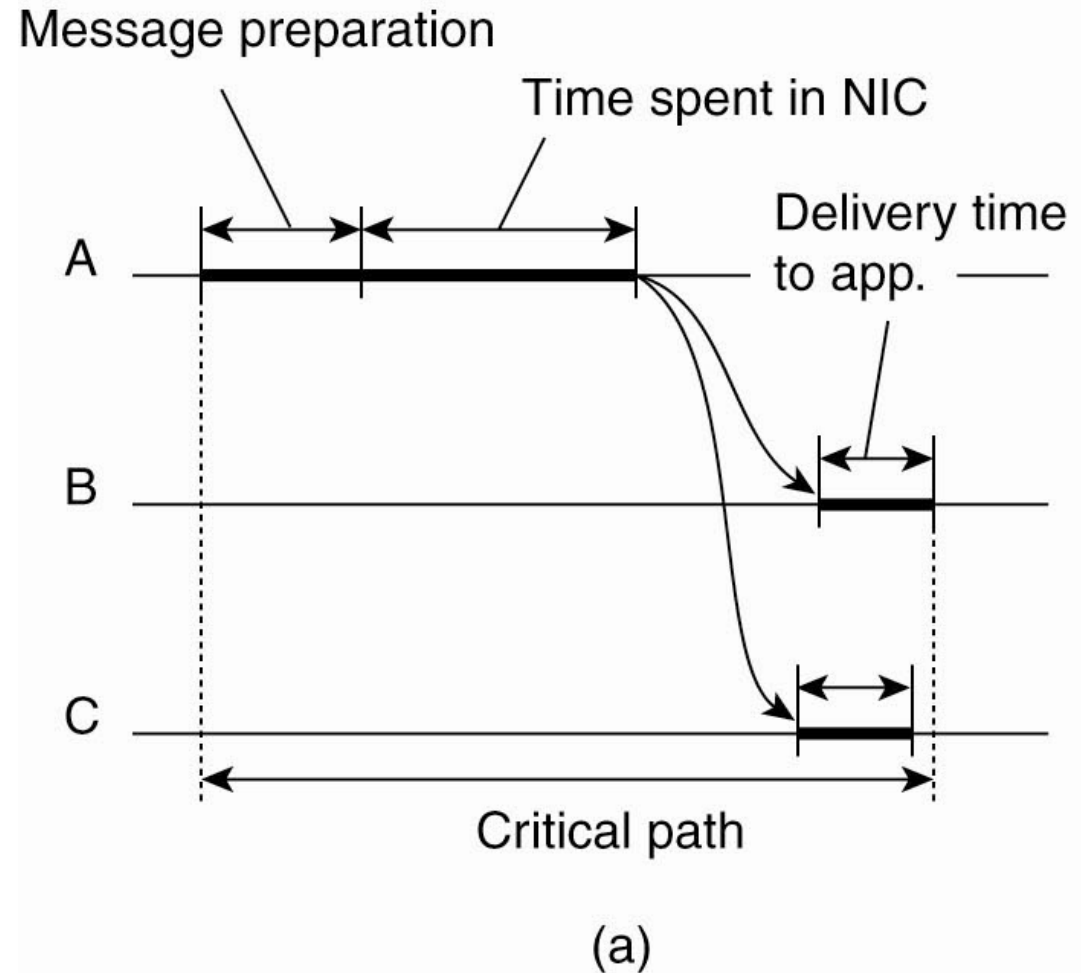
- The time daemon tells everyone how to adjust their clock.





# Clock Synchronization in Wireless Networks (1)

- The usual critical path in determining network delays.

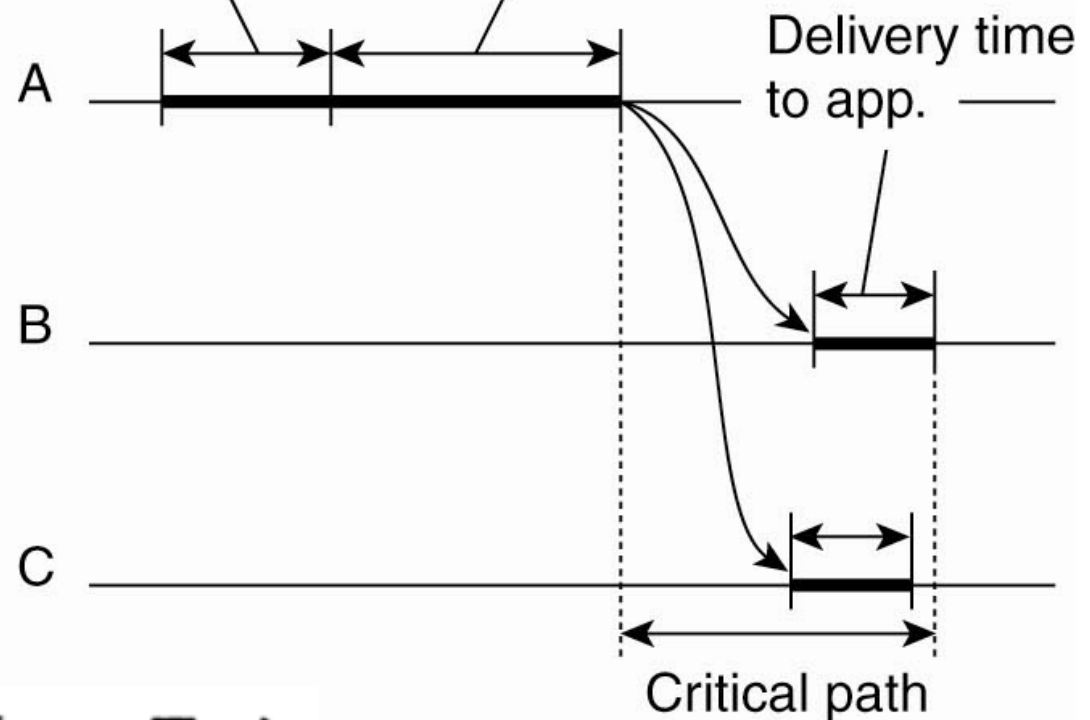


# Clock Synchronization in Wireless Networks (2)

- The critical path in the case of RBS.

Message preparation

Time spent in NIC



$$Offset [p, q] = \frac{\sum_{k=1}^M (T_{p,k} - T_{q,k})}{M}$$

(b)

## 2. Logical clock



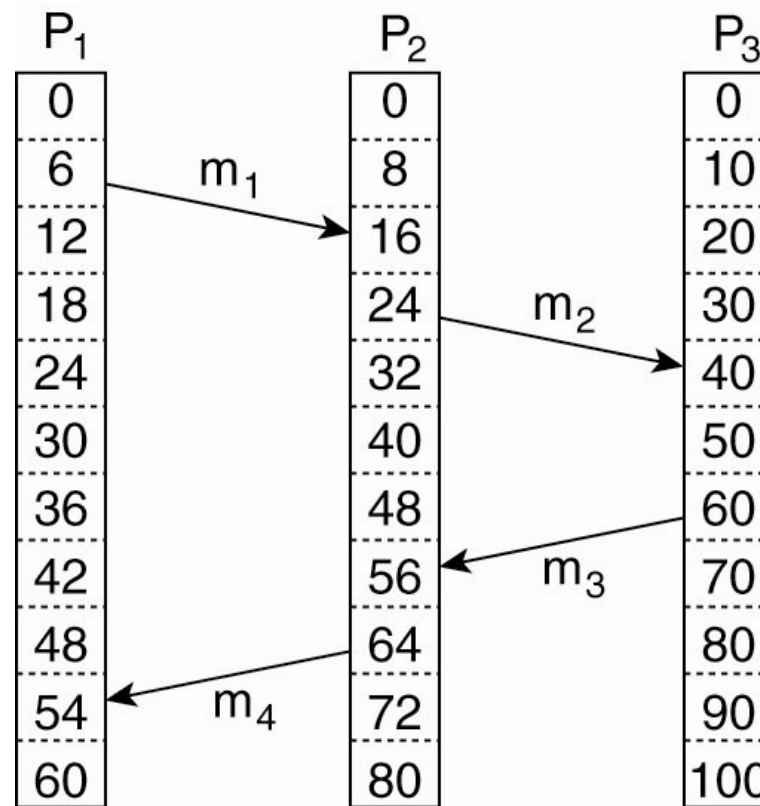
- Lamport logical clocks
- Vector clocks

## 2.1. Lamport's Logical Clocks (1)

- The "happens-before" relation  $\rightarrow$  can be observed directly in two situations:
  1. If  $a$  and  $b$  are events in the same process, and  $a$  occurs before  $b$ , then  $a \rightarrow b$  is true.
  2. If  $a$  is the event of a message being sent by one process, and  $b$  is the event of the message being received by another process, then  $a \rightarrow b$
- $a \rightarrow b$  and  $b \rightarrow c$ , then  $a \rightarrow c$
- Concurrent

# Lamport's Logical Clocks (2)

- Three processes, each with its own clock. The clocks run at different rates.



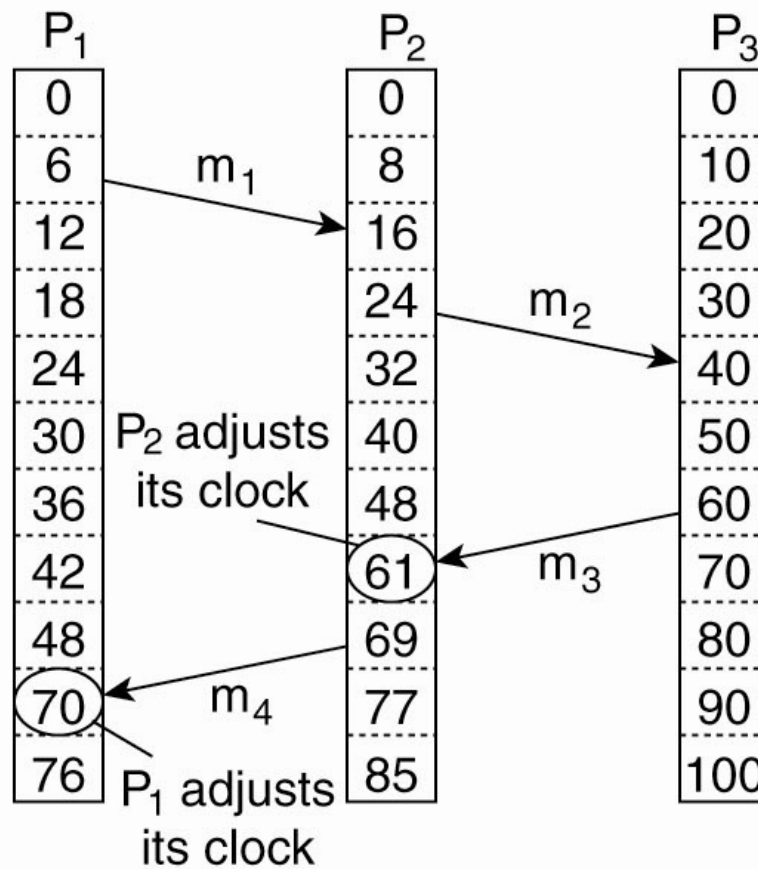
(a)

# Lamport's Logical Clocks (5)

- Updating counter  $C_i$  for process  $P_i$ 
  1. Before executing an event  $P_i$  executes
$$C_i \leftarrow C_i + 1.$$
  2. When process  $P_i$  sends a message  $m$  to  $P_j$ , it sets  $m$ 's timestamp  $ts(m)$  equal to  $C_i$  after having executed the previous step.
  3. Upon the receipt of a message  $m$ , process  $P_j$  adjusts its own local counter as
$$C_j \leftarrow \max\{C_j, ts(m)\},$$
 after which it then executes the first step and delivers the message to the application.

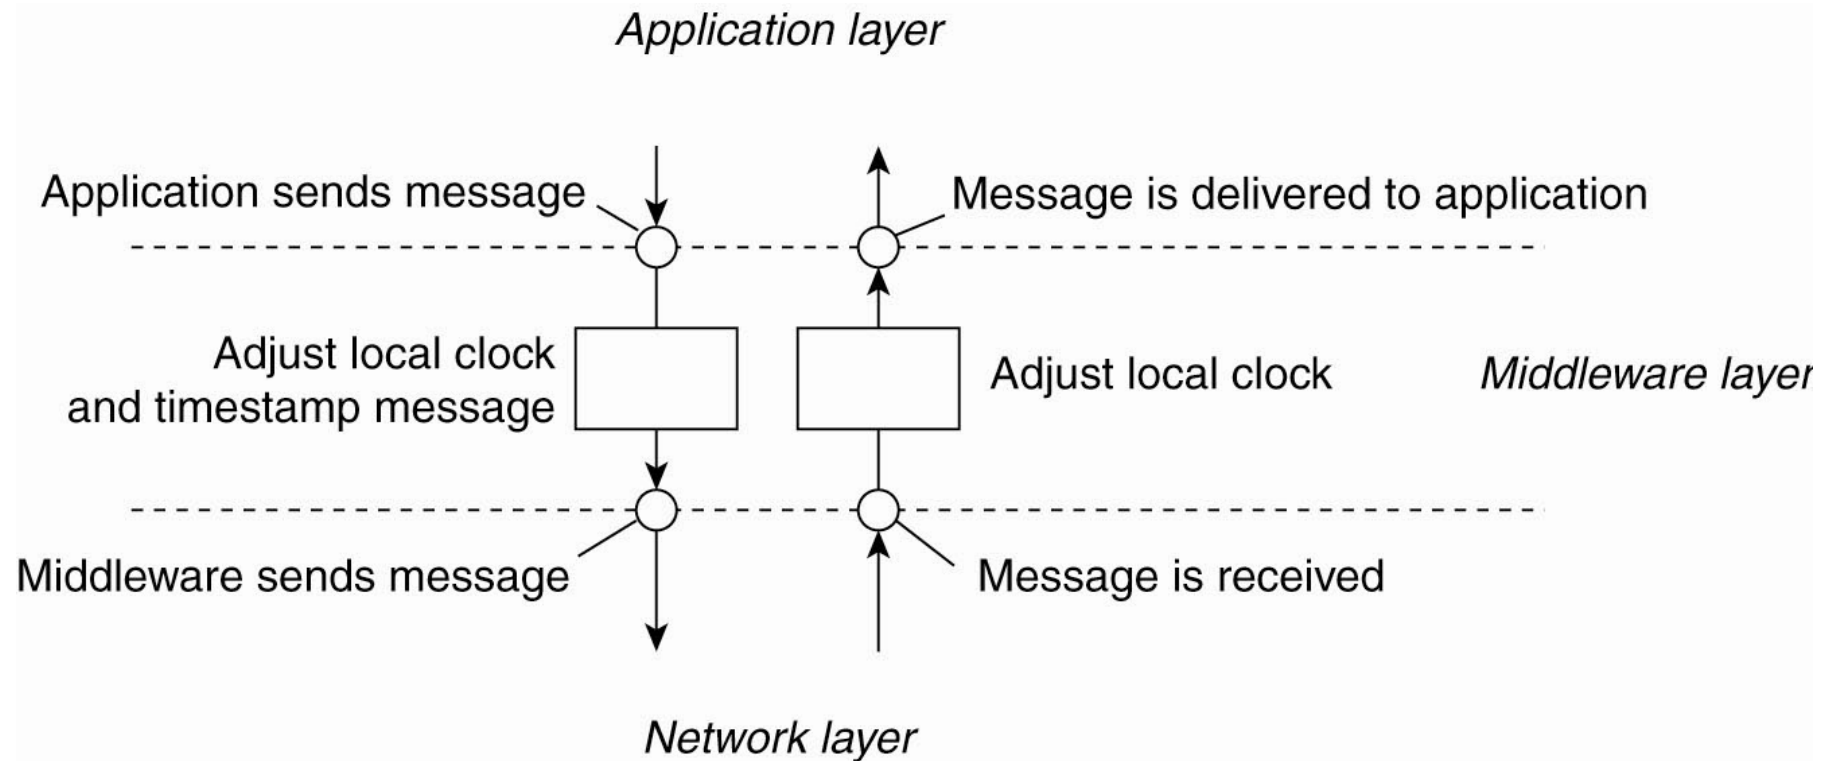
# Lamport's Logical Clocks (3)

(b) Lamport's algorithm corrects the clocks.



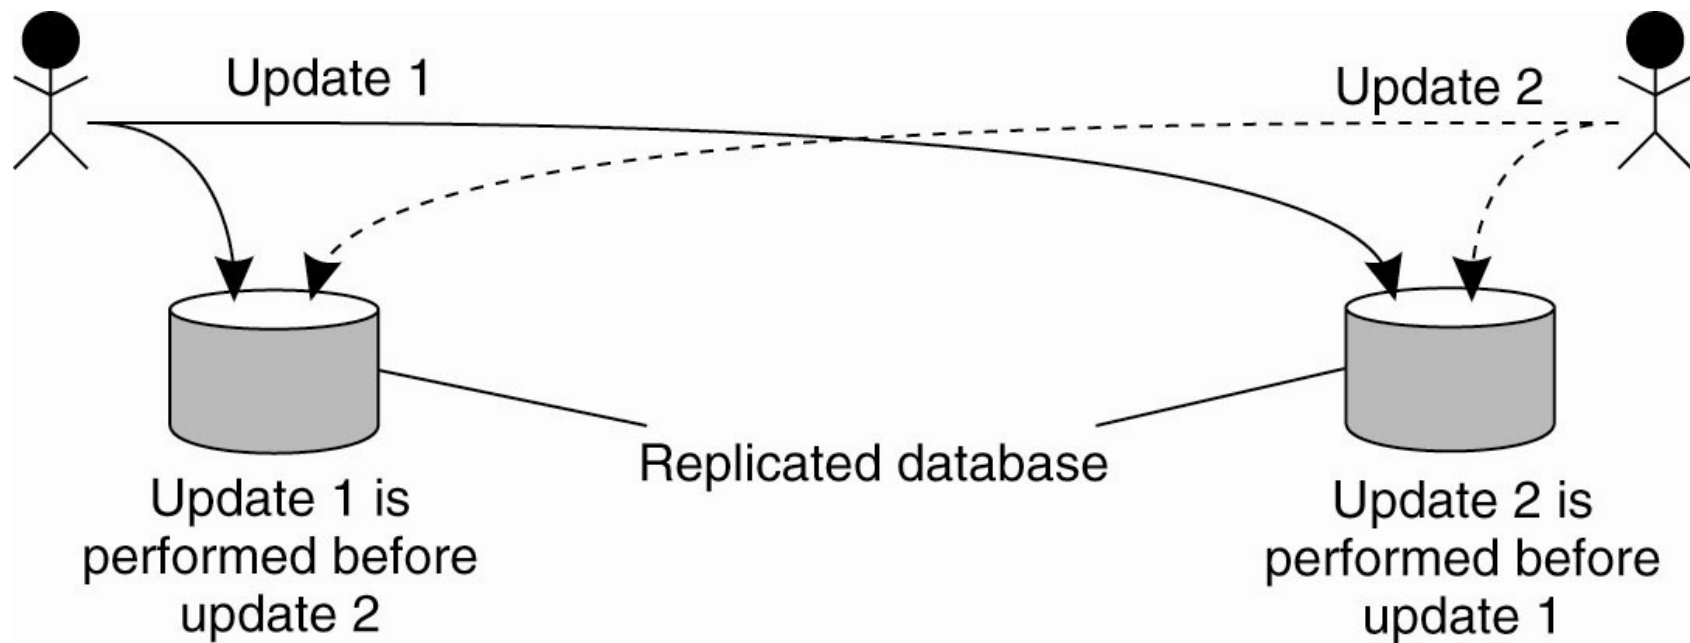
(b)

# Lamport's Logical Clocks (4)





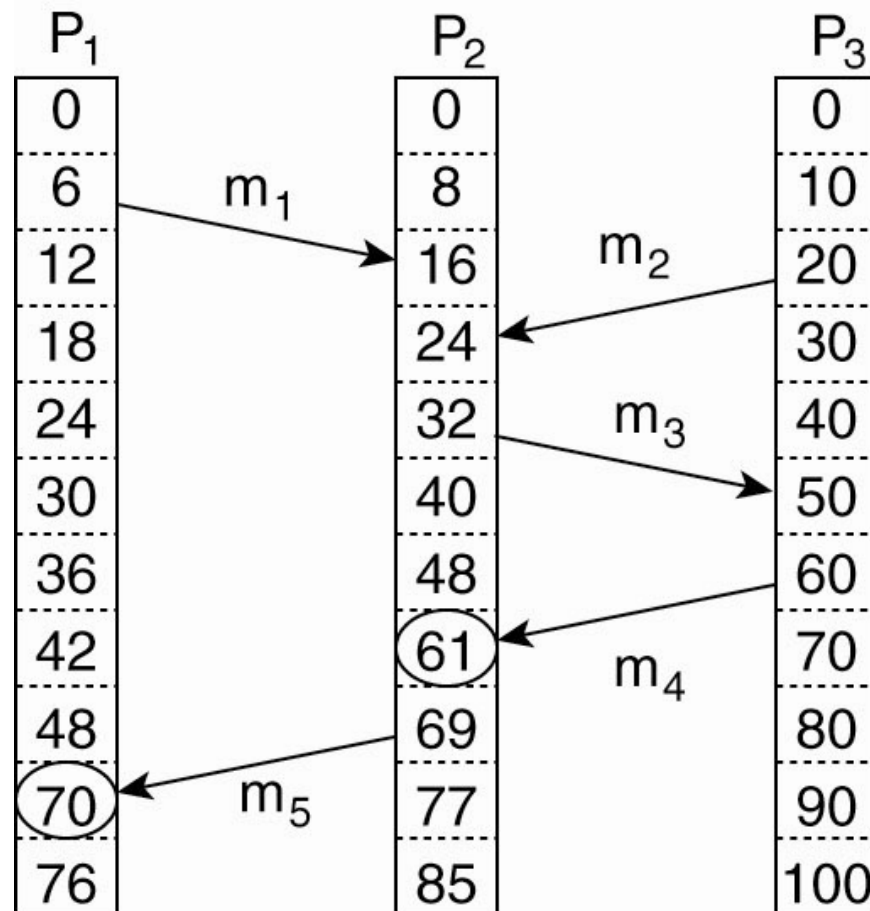
# Example: Totally Ordered Multicasting



Updating a replicated database and leaving it in an inconsistent state.

## 2.2. Vector Clocks (1)

- Concurrent message transmission using logical clocks.



# Vector Clocks (2)

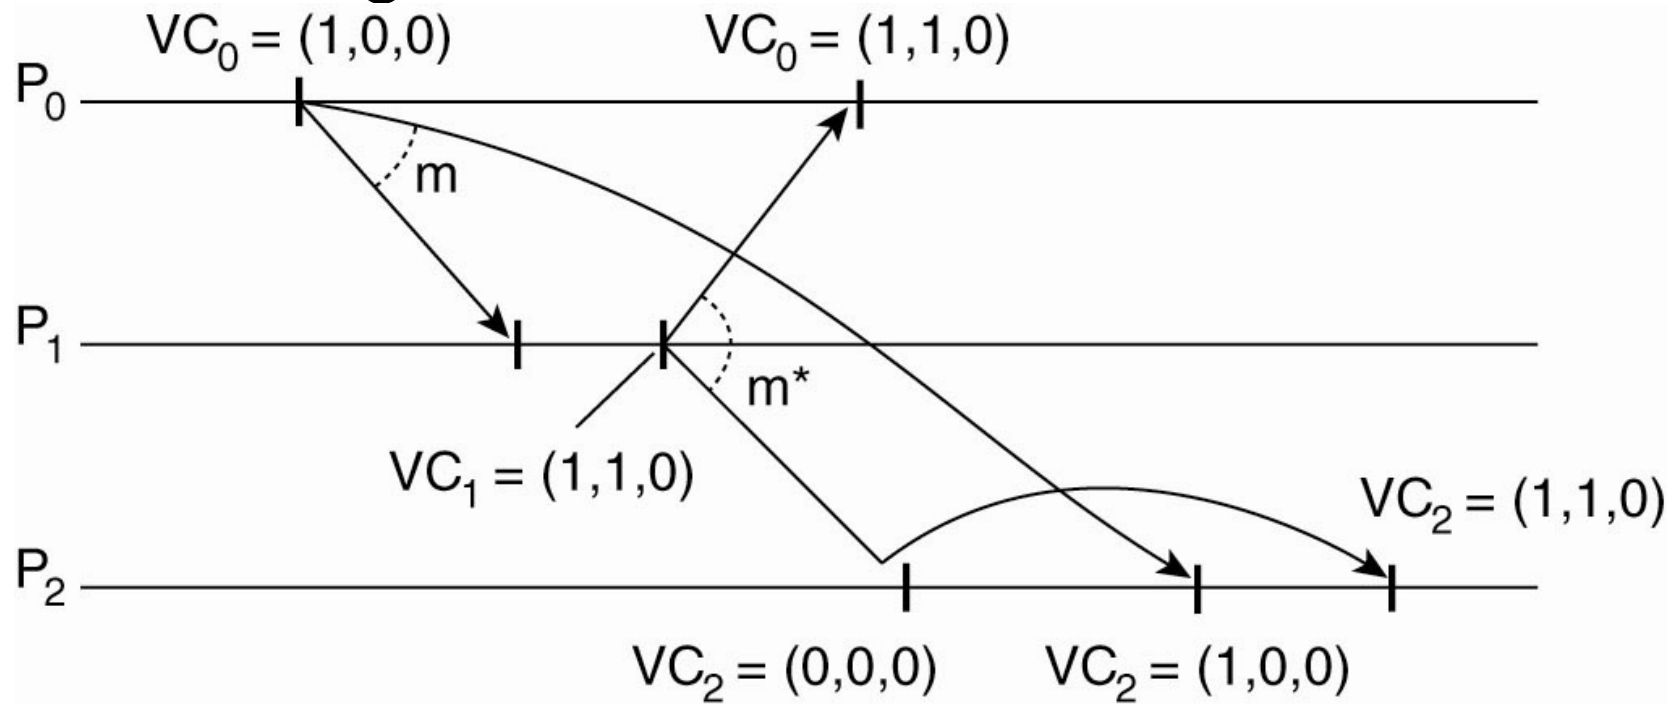
- Vector clocks are constructed by letting each process  $P_i$  maintain a vector  $VC_i$  with the following two properties:
  1.  $VC_i[i]$  is the number of events that have occurred so far at  $P_i$ . In other words,  $VC_i[i]$  is the local logical clock at process  $P_i$ .
  2. If  $VC_i[j] = k$  then  $P_i$  knows that  $k$  events have occurred at  $P_j$ . It is thus  $P_i$ 's knowledge of the local time at  $P_j$ .

# Vector Clocks (3)

- Steps carried out to accomplish property 2 of previous slide:
  1. Before executing an event  $P_i$  executes  $VC_i[i] \leftarrow VC_i[i] + 1$ .
  2. When process  $P_i$  sends a message  $m$  to  $P_j$ , it sets  $m$ 's (vector) timestamp  $ts(m)$  equal to  $VC_i$  after having executed the previous step.
  3. Upon the receipt of a message  $m$ , process  $P_j$  adjusts its own vector by setting  $VC_j[k] \leftarrow \max\{VC_j[k], ts(m)[k]\}$  for each  $k$ , after which it executes the first step and delivers the message to the application.

# Enforcing Causal Communication

□ Enforcing causal communication.



2 conditions:

1.  $ts(m)[i] = VC_j[i] + 1$
2.  $ts(m)[k] \leq VC_j[k]$  for all  $k \neq i$

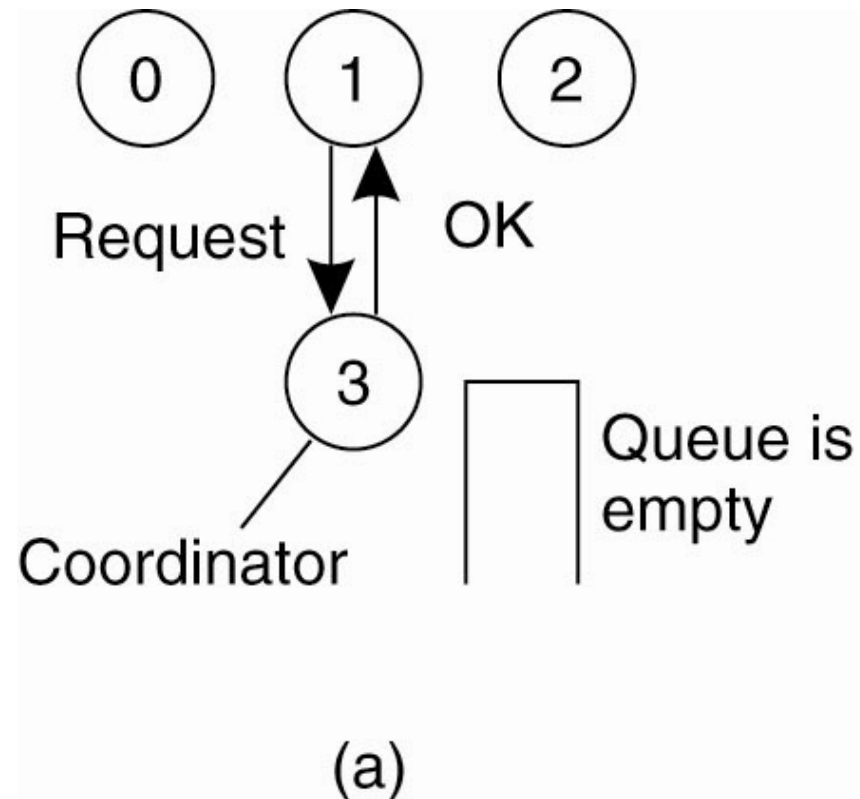
# 3. Mutual exclusion



- ❑ A Centralized Algorithm
- ❑ A Decentralized Algorithm
- ❑ A Distributed Algorithm
- ❑ A Token Ring Algorithm
- ❑ A Comparison of the Three Algorithms

# Mutual Exclusion

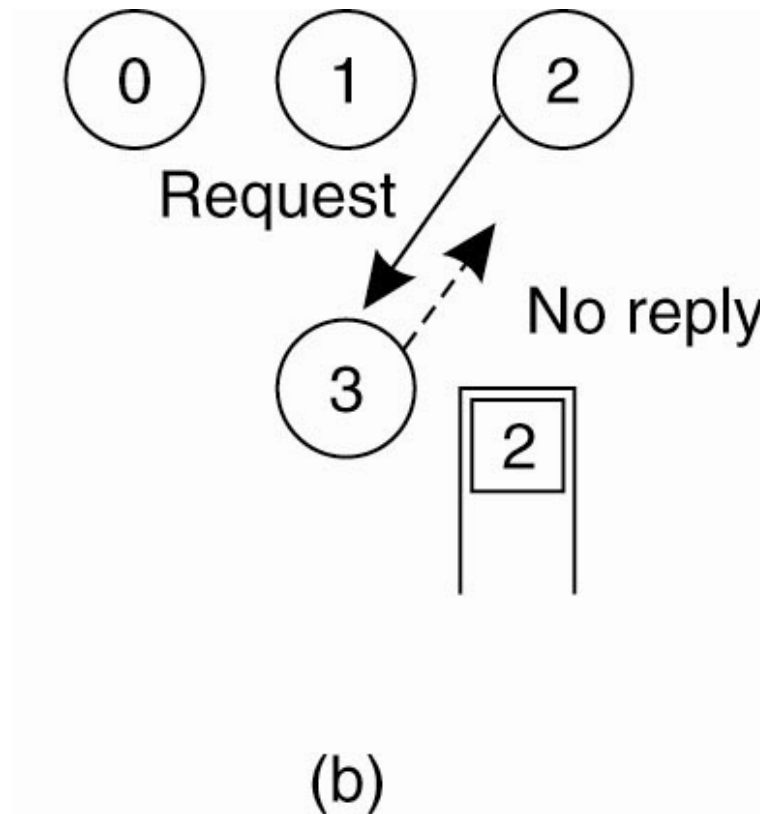
## 3.1. A Centralized Algorithm (1)



- Process 1 asks the coordinator for permission to access a shared resource. Permission is granted.

# Mutual Exclusion

## A Centralized Algorithm (2)



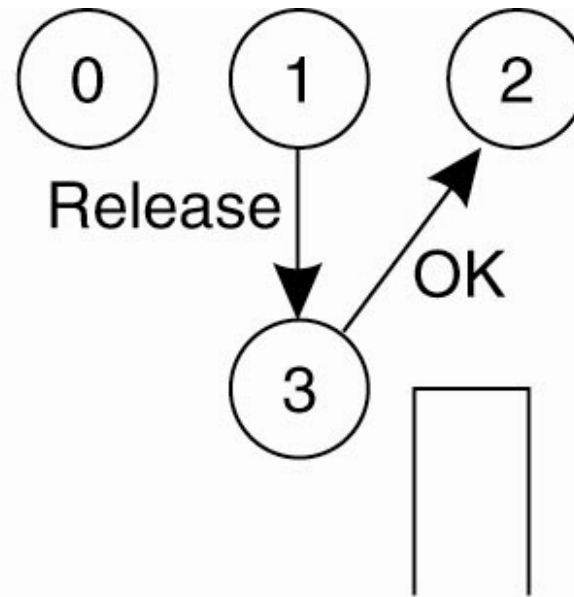
- Process 2 then asks permission to access the same resource. The coordinator does not reply.



# Mutual Exclusion

## A Centralized Algorithm (3)

- When process 1 releases the resource, it tells the coordinator, which then replies to 2.



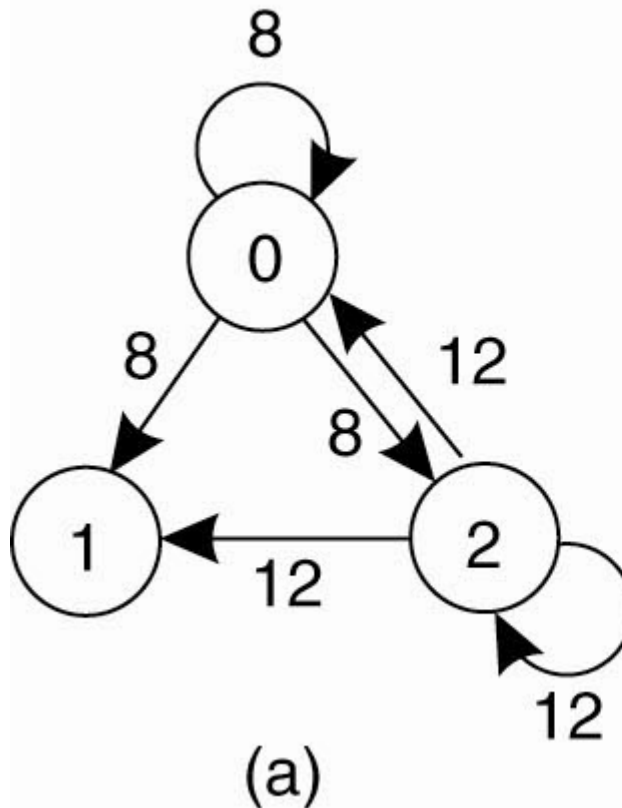
(c)

## 3.2. A Distributed Algorithm (1)

- Three different cases:
  1. If the receiver is not accessing the resource and does not want to access it, it sends back an OK message to the sender.
  2. If the receiver already has access to the resource, it simply does not reply. Instead, it queues the request.
  3. If the receiver wants to access the resource as well but has not yet done so, it compares the timestamp of the incoming message with the one contained in the message that it has sent everyone. The lowest one wins.

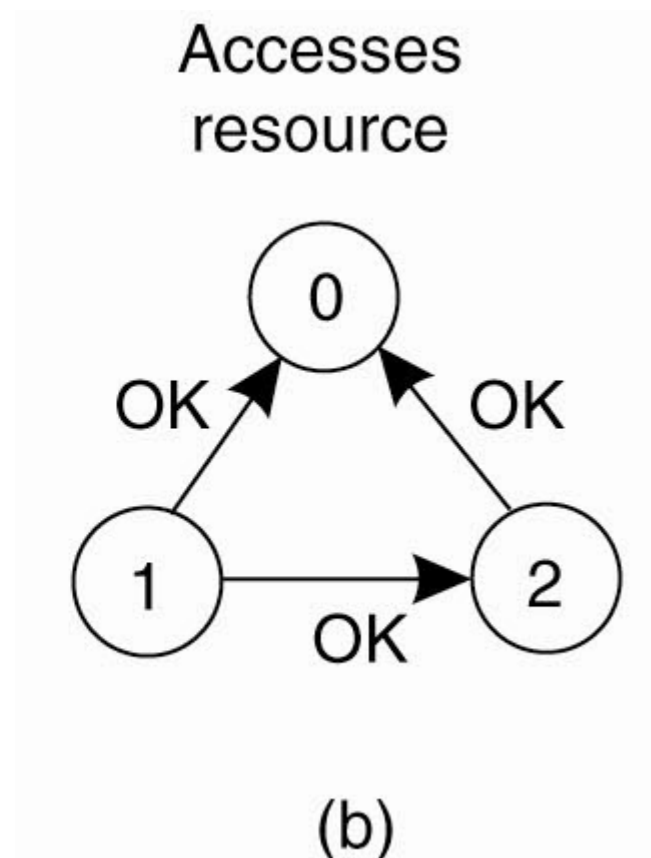
# A Distributed Algorithm (2)

- Two processes want to access a shared resource at the same moment.



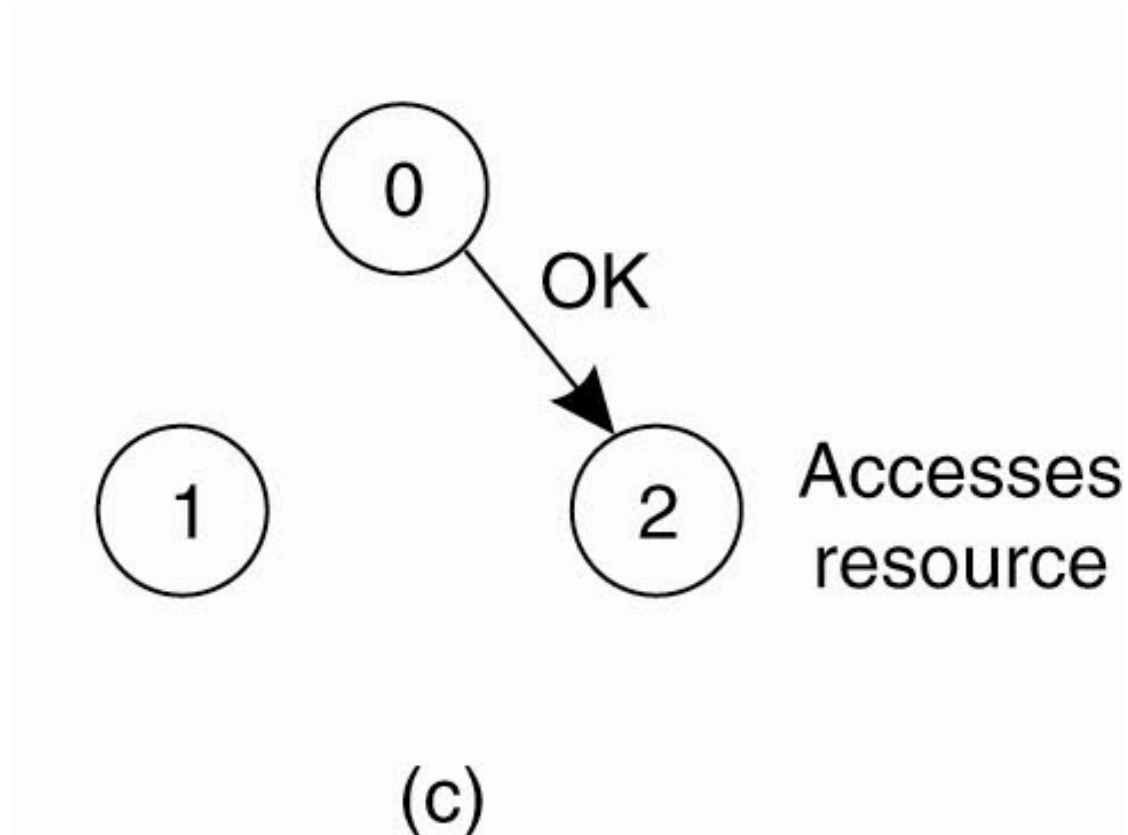
# A Distributed Algorithm (3)

- Process 0 has the lowest timestamp, so it wins.

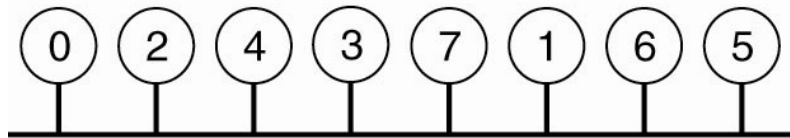


# A Distributed Algorithm (4)

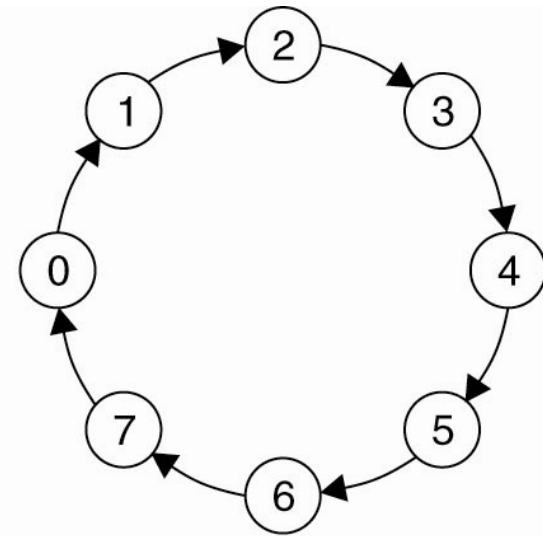
- When process 0 is done, it sends an OK also, so 2 can now go ahead.



## 3.3. A Token Ring Algorithm



(a)

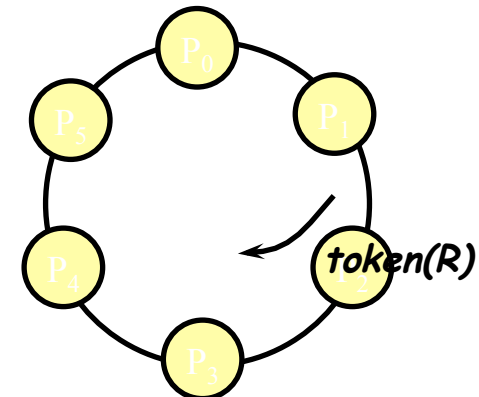


(b)

- (a) An unordered group of processes on a network.
- (b) A logical ring constructed in software.

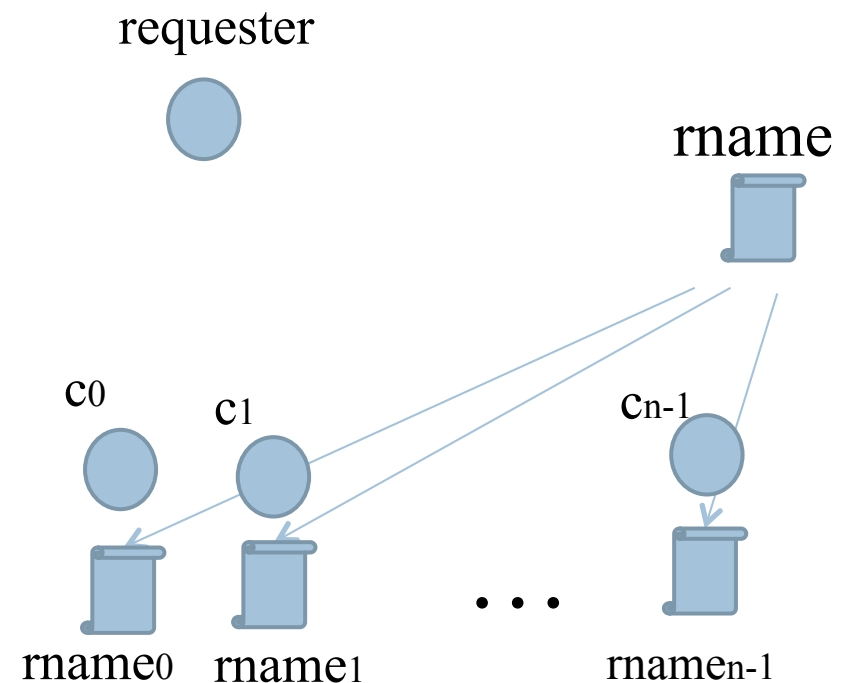
# Token Ring algorithm

- Initialization
  - ▣ Process 0 gets token for resource R
- Token circulates around ring
  - ▣ From  $P_i$  to  $P_{(i+1) \bmod N}$
- When process acquires token
  - ▣ Checks to see if it needs to enter critical section
  - ▣ If no, send token to neighbor
  - ▣ If yes, access resource
    - Hold token until done



## 3.4. Decentralized Algorithm

- Based on the Distributed Hash Table (DHT) system structure previously introduced
  - ▣ Peer-to-peer
  - ▣ Object names are hashed to find the successor node that will store them
- Here, we assume that  $n$  replicas of each object are stored





# Placing the Replicas

- The resource is known by a unique name: *rname*
  - ▣ Replicas: *rname-0*, *rname-1*, ..., *rname-(n-1)*
  - ▣ *rname-i* is stored at *succ(rname-i)*, where names and site names are hashed as before
  - ▣ If a process knows the name of the resource it wishes to access, it also can generate the hash keys that are used to locate all the replicas

# The Decentralized Algorithm



- Every replica has a coordinator that controls access to it (the coordinator is the node that stores it)
- For a process to use the resource it must receive permission from  $m > n/2$  coordinators
- This guarantees exclusive access as long as a coordinator only grants access to one process at a time

# The Decentralized Algorithm



- The coordinator notifies the requester when it has been denied access as well as when it is granted
  - ▣ Requester must “count the votes”, and decide whether or not overall permission has been granted or denied
- If a process (requester) gets fewer than  $m$  votes it will wait for a random time and then ask again

# 4. Election Algorithms



- Traditional Election algorithms
  - ▣ The Bully Algorithm
  - ▣ A Ring Algorithm
- Election in Wireless Environments
- Election in Large-Scale Systems

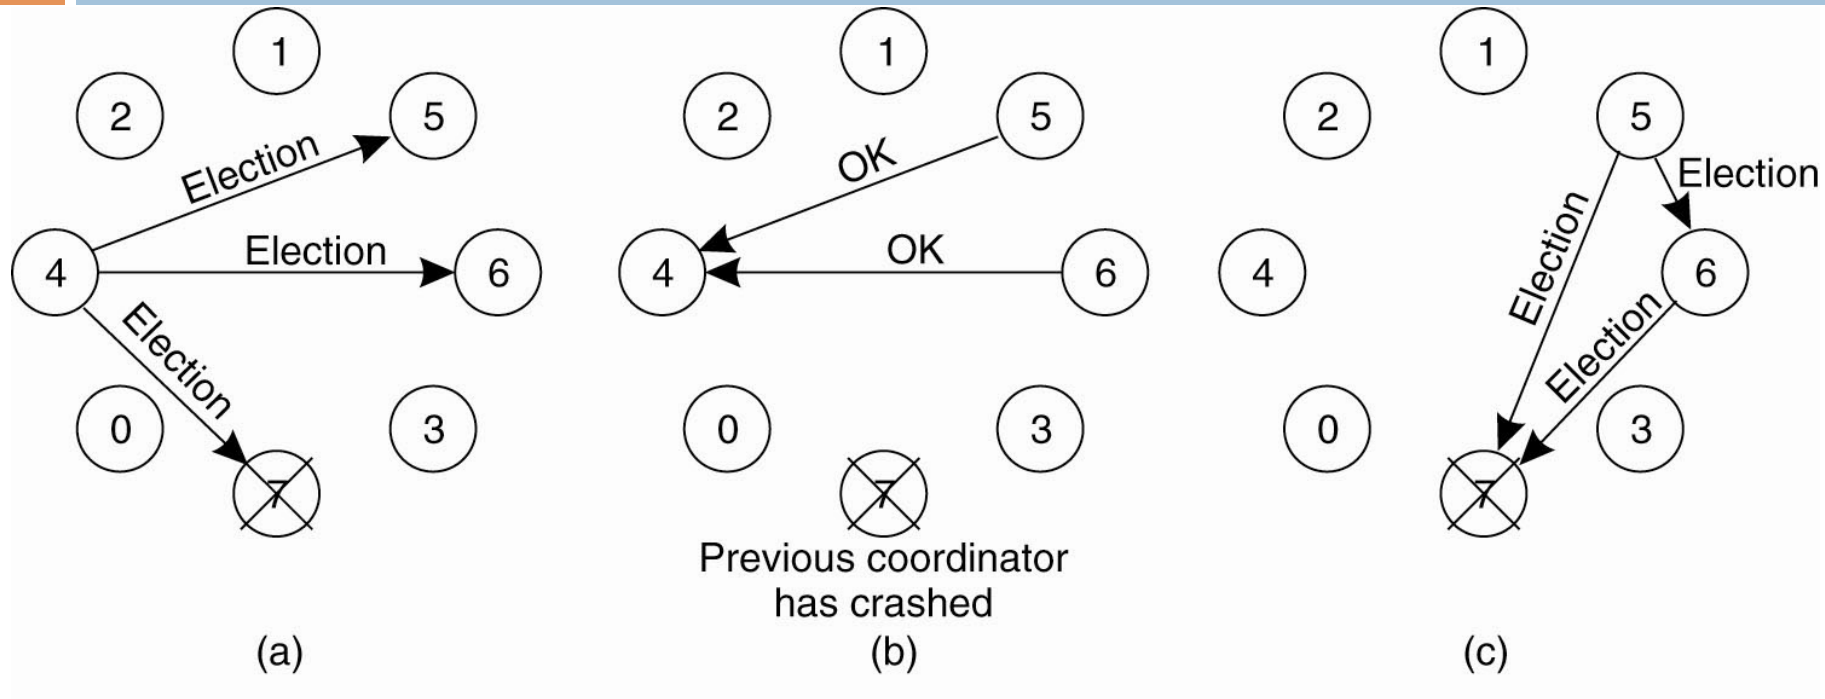
# Election Algorithms



## □ The Bully Algorithm

1.  $P$  sends an *ELECTION* message to all processes with higher numbers.
2. If no one responds,  $P$  wins the election and becomes coordinator.
3. If one of the higher-ups answers, it takes over.  $P$ 's job is done.

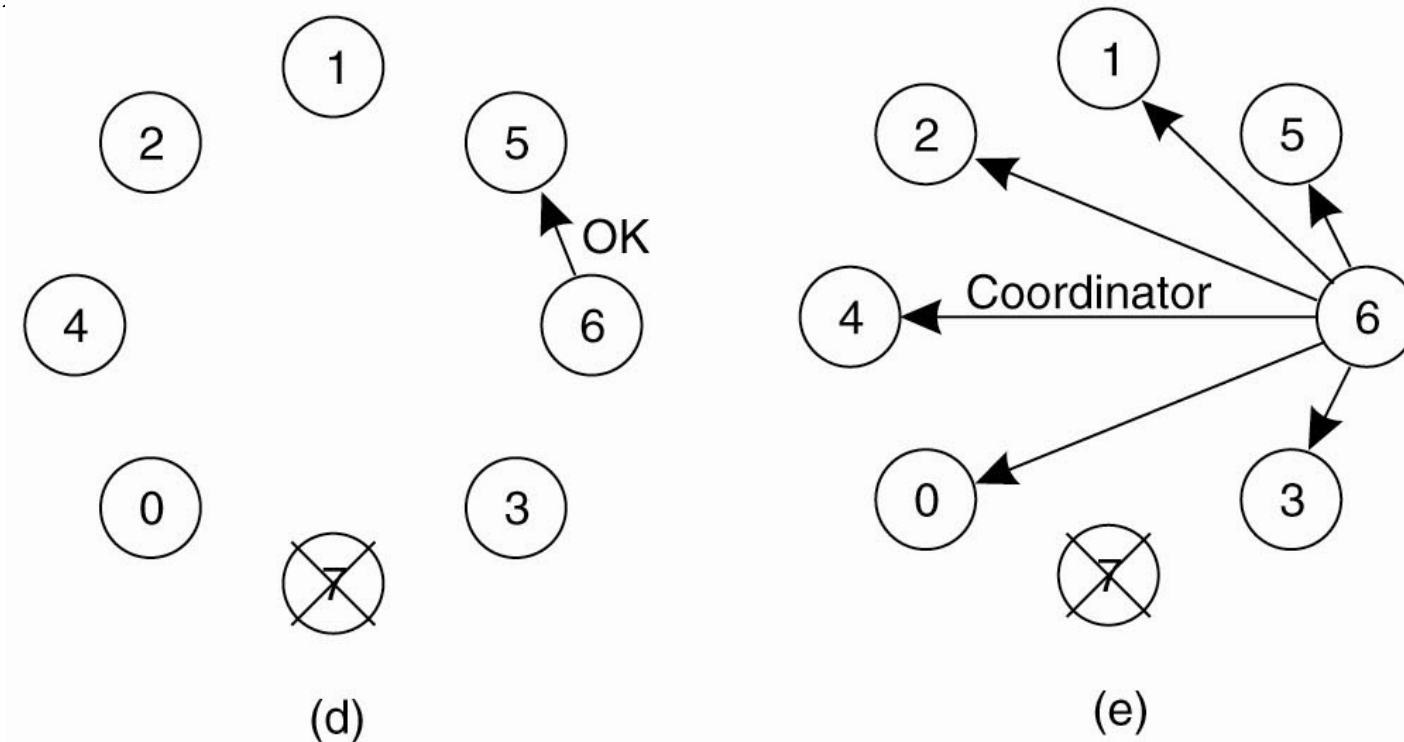
# The Bully Algorithm (1)



- The bully election algorithm. (a) Process 4 holds an
- election. (b) Processes 5 and 6 respond, telling 4 to stop.
- (c) Now 5 and 6 each hold an election.

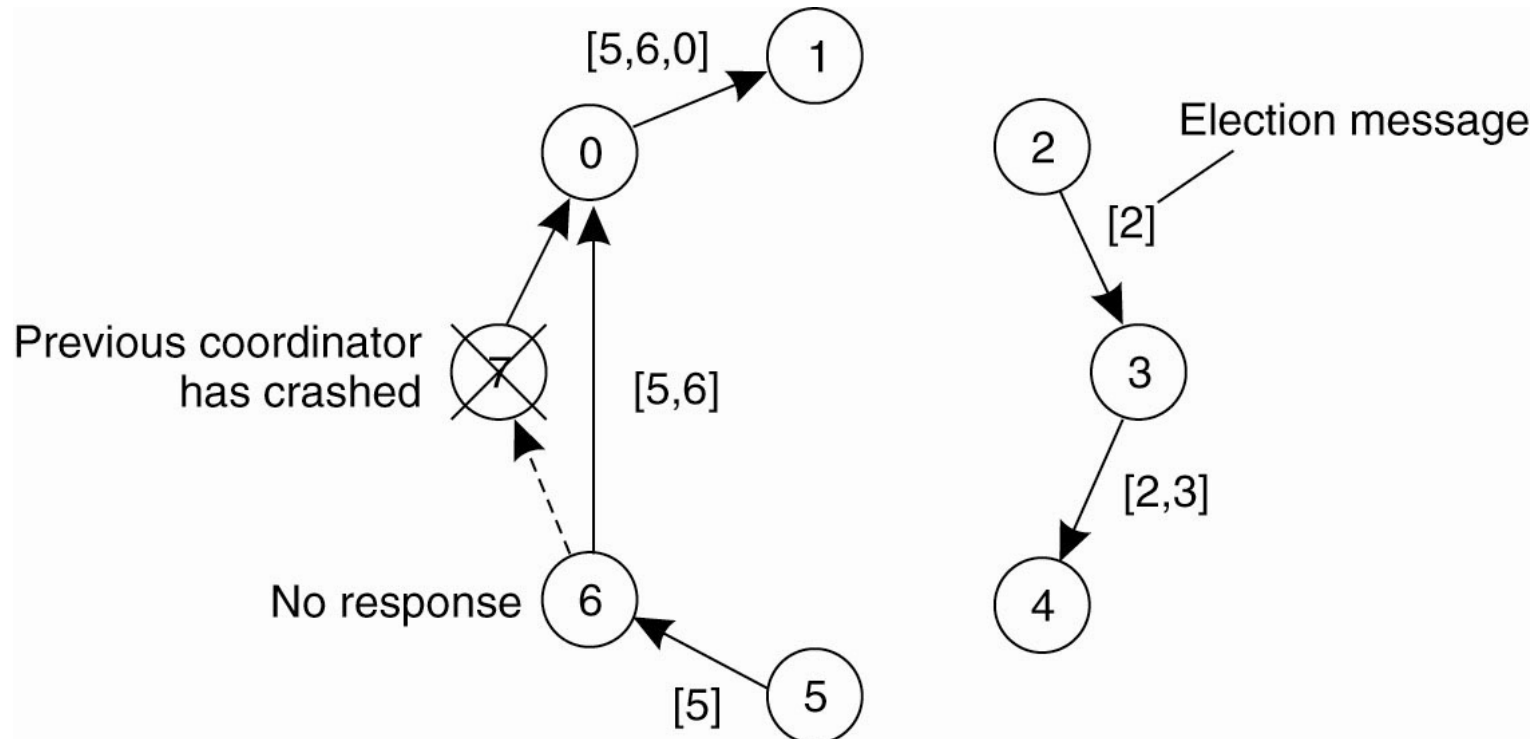
# The Bully Algorithm (2)

- The bully election algorithm. (d) Process 6 tells 5



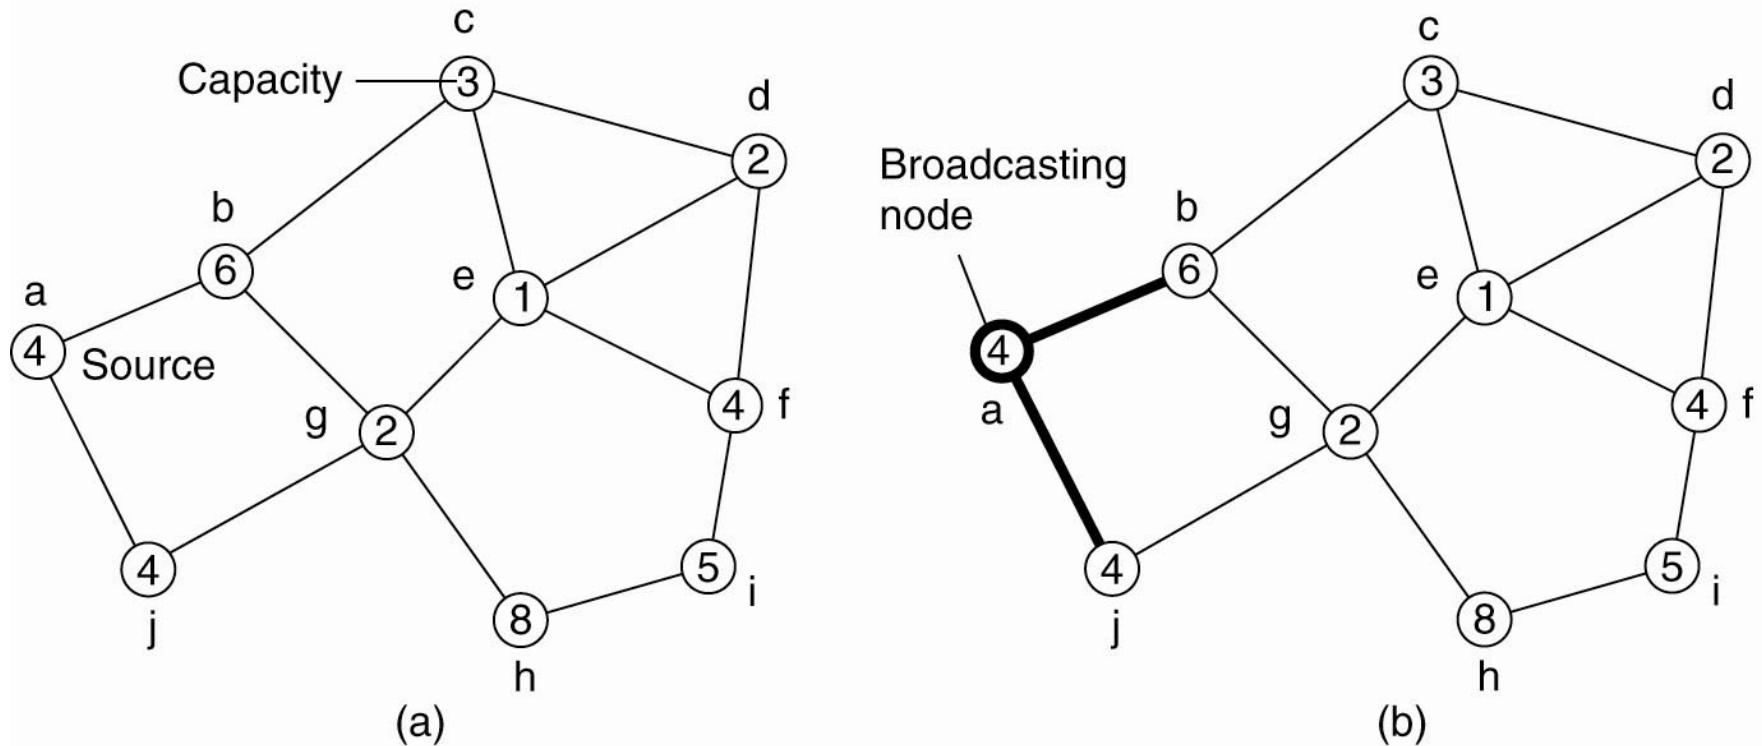
# A Ring Algorithm

- Election algorithm using a ring.



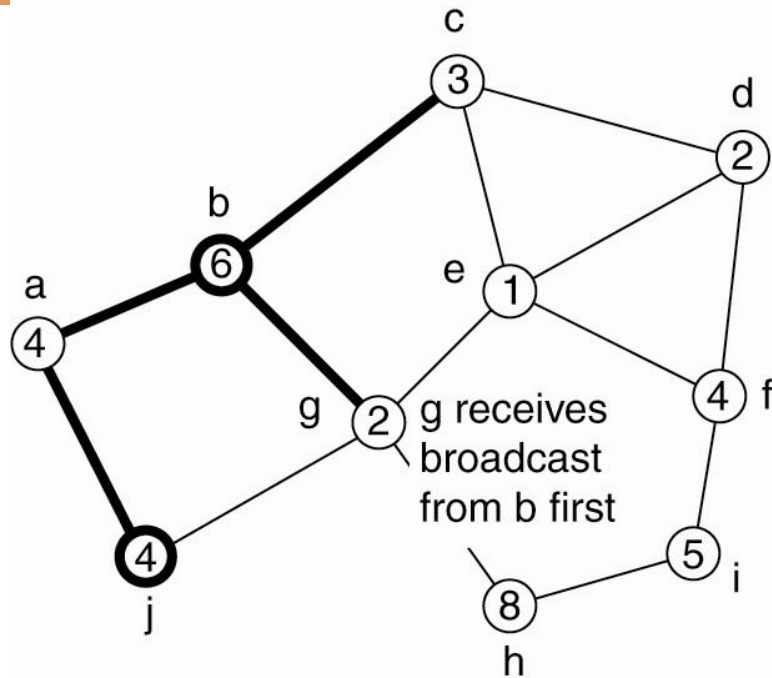


# Elections in Wireless Environments (1)

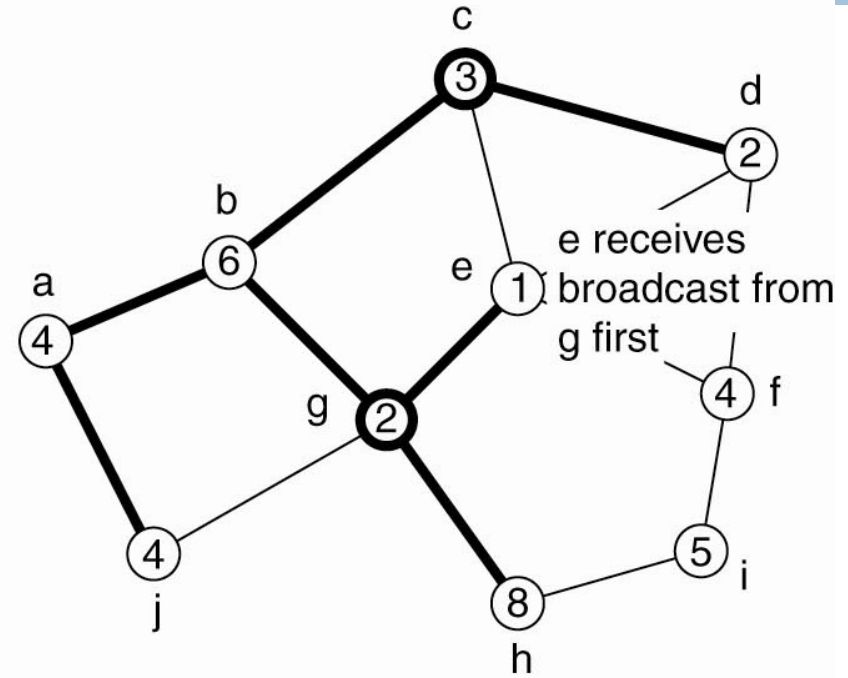


- Election algorithm in a wireless network, with node a as the source. (a) Initial network. (b)–(e) The build-tree phase

# Elections in Wireless Environments (2)



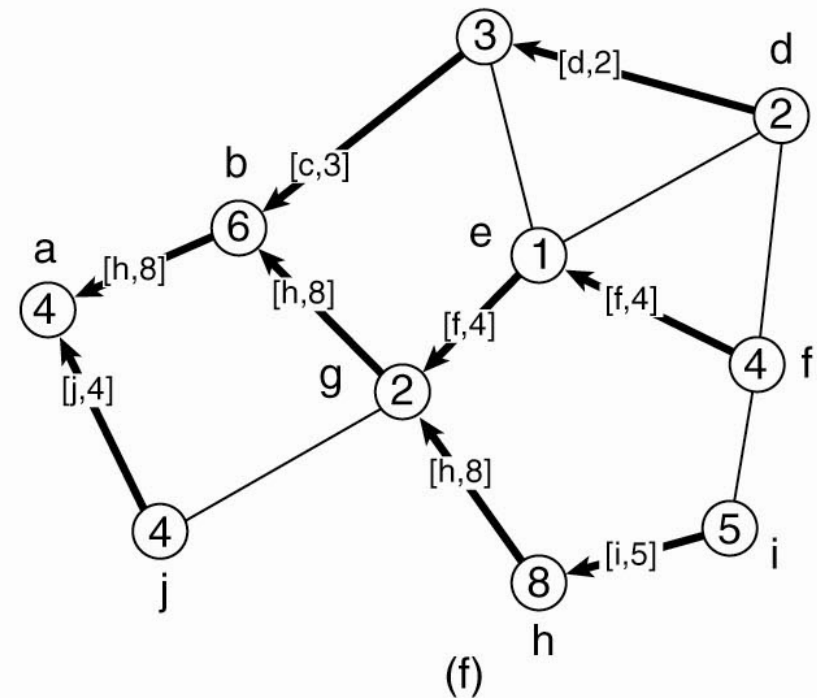
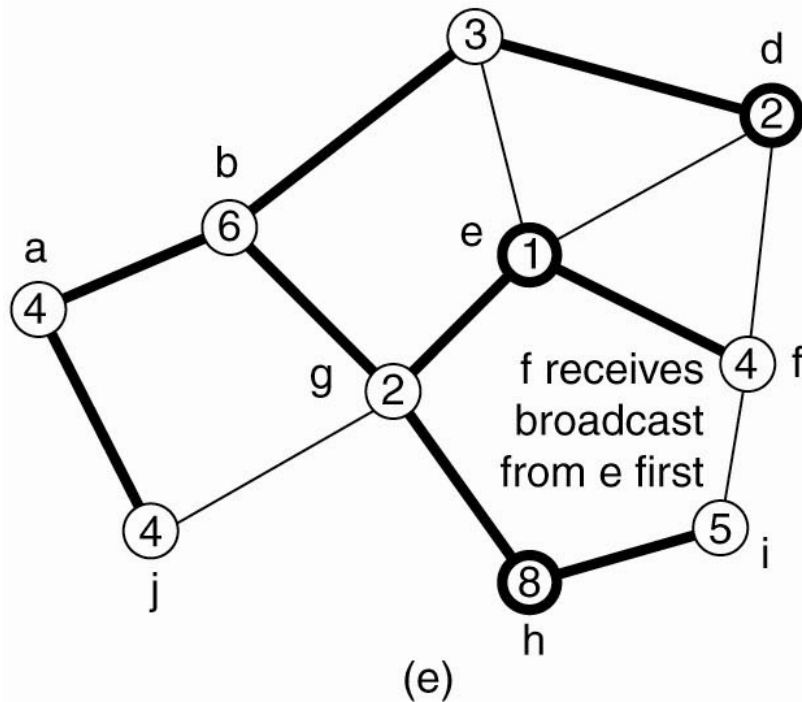
(c)



(d)

# Elections in Wireless Environments (3)

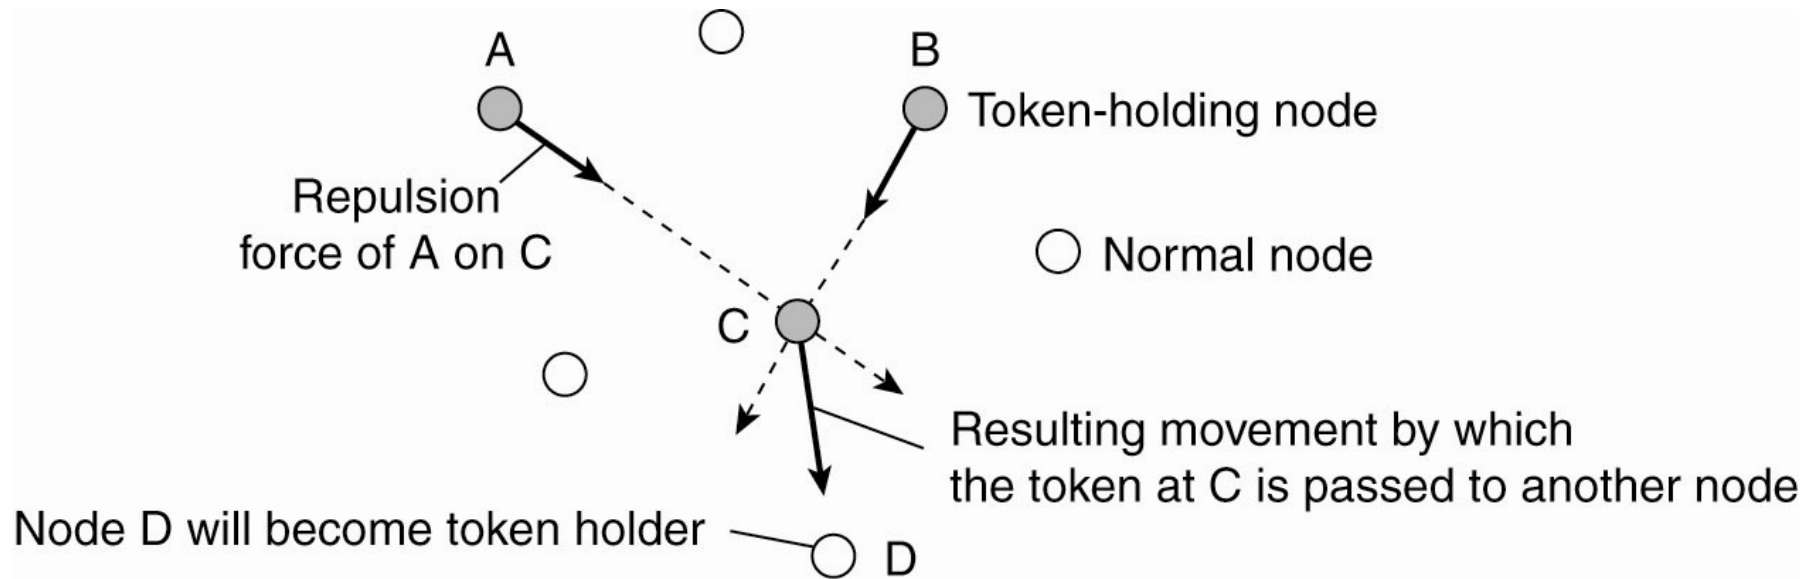
- (e) The build-tree phase.
- (f) Reporting of best node to source.



# Elections in Large-Scale Systems (1)

- Requirements for superpeer selection:
  1. Normal nodes should have low-latency access to superpeers.
  2. Superpeers should be evenly distributed across the overlay network.
  3. There should be a predefined portion of superpeers relative to the total number of nodes in the overlay network.
  4. Each superpeer should not need to serve more than a fixed number of normal nodes.

# Elections in Large-Scale Systems (2)



- Moving tokens in a two-dimensional space using repulsion forces.